

# Denial of Service

Computer Security 2 (GA02)

Emiliano De Cristofaro, Gianluca Stringhini

Thanks to Aurelien Francillon for  
letting us re-use some of the slides

# Reading

Stallings 21.5

# What?

- Denial of Service, in short DoS
- Goal (informal): to take out a large site with little computing work
- How: Amplification
  - Use small number of packets  $\Rightarrow$  obtain a big effect
- (Roughly speaking) two types of amplification attacks:
  - DoS bug:
    - Design flaw allowing one machine to disrupt a service
  - DoS flood:
    - Command botnet to generate flood of requests

# A high profile example: Estonia



- Attacked sites: (started April 2007, lasted two weeks)
  - Estonian ministerial sites
  - Various Estonian commercial sites

# DoS can happen at any layer


- In this lecture, you'll learn about
  - DoS at different **layers**
  - Generic DoS solutions
  - Network DoS solutions
- Sad truth:
  - Internet **not designed** to handle DoS and Distributed DoS (DDoS) attacks

(Challenges → Opportunities 😊)

# 802.11b DoS bugs

- Wireless radio **jamming** attacks:
  - Trivial, not our focus.
- Protocol DoS bugs: [Bellardo, Savage, 2003]
  - **NAV** (Network Allocation Vector):
    - 15-bit field, max value  $2^{15}-1=32767$
    - **Any node** can reserve channel for **#NAV seconds**
    - No one else should transmit during NAV period
  - **De-authentication** bug:
    - Any node can send de-auth packet to AP
    - Deauth packet unauthenticated
    - ... attacker can repeatedly deauth (i.e., kick out) anyone


# Smurf amplification DoS attack



- Send ping request to broadcast addr (ICMP Echo Req)
- Lots of responses:
  - Every host on target network generates a ping reply (ICMP Echo Reply) to victim
- Prevention: reject external packets to broadcast address

# DNS Amplification Attacks

×50 amplification



(\*) EDNS -> Extension mechanisms for DNS, i.e., a way to add parameters to DNS messages (due to prior size restrictions)

2006: 0.58M open resolvers on Internet (Kaminsky-Shiffman)

2013: 21.7M open resolvers (openresolverproject.org)

⇒ 3/2013: DDoS attack generating 300 Gbps



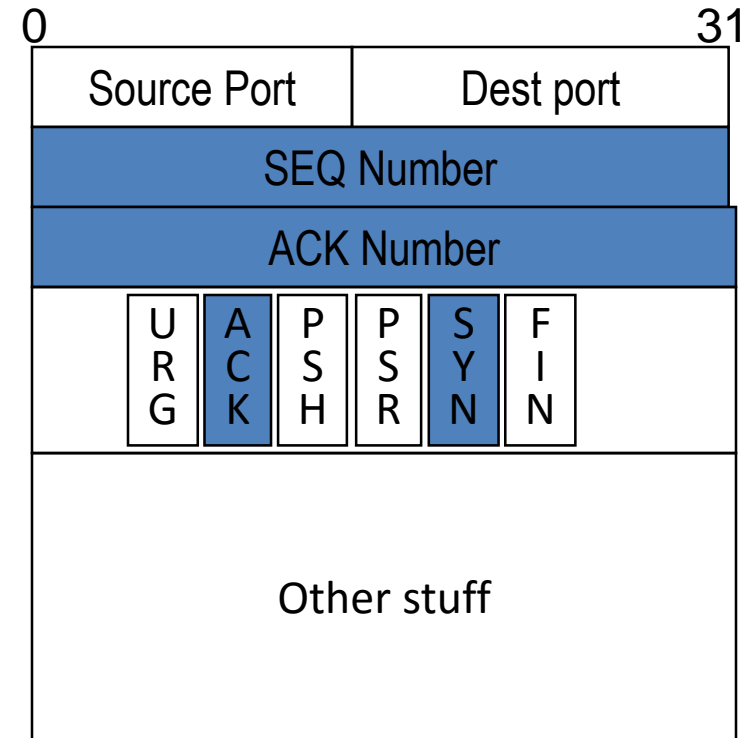
# Review: IP Header format

- Connectionless
  - Unreliable
  - Best effort


0	31
Version	Header Length
Type of Service	
Total Length	
Identification	
Flags	Fragment Offset
Time to Live	
Protocol	
Header Checksum	
Source Address of Originating Host	
Destination Address of Target Host	
Options	
Padding	
IP Data	

# Review: TCP Header format


- TCP:
  - Session based
  - Congestion control
  - In order delivery



# Review: TCP Handshake



# TCP SYN Flood I: low rate (DoS bug)



## Single machine:

- SYN Packets with **random source IP addresses**
- Fills up backlog queue on server
- No further connections possible

# SYN Floods

OS	Backlog queue size
<b>Linux 1.2.x</b>	10
<b>FreeBSD 2.1.5</b>	128
<b>WinNT 4.0</b>	6

Backlog timeout: 3 minutes

Attacker needs only send 128 SYN packets every 3 minutes.

Low rate SYN flood

# A classic SYN flood example

- MS Blaster worm (2003)
  - Infected machines at noon on Aug 16<sup>th</sup>:
    - SYN flood on port 80 to **windowsupdate.com**
    - 50 SYN packets every second.
      - each packet is 40 bytes.
    - Spoofed source IP: a.b.X.Y where X,Y random.
- MS solution:
  - new name: **windowsupdate.microsoft.com**
  - Win update file delivered by Akamai

# Low rate SYN flood defenses

- Non-solution:
  - Increase backlog queue size or decrease timeout
- **Correct solution** (when under attack) :
  - **Syncookies**: remove state from server
  - Small performance overhead

# Syncookies

- Use secret key and data in packet to gen. server SN
- Server responds to Client with SYN-ACK cookie:
  - $T = 5\text{-bit counter incremented every 64 secs.}$
  - $L = \text{MAC}_{\text{key}} (\text{SAddr}, \text{SPort}, \text{DAddr}, \text{DPort}, \text{SN}_C, T)$  [24 bits]
    - key: picked at random during boot
  - $\text{SN}_S = (T \cdot \text{mss} \cdot L)$  (  $|L| = 24 \text{ bits}$  )
  - **Server does not save state** (other TCP options are lost)
- Honest client responds with ACK (  $\text{AN}=\text{SN}_S$  ,  $\text{SN}=\text{SN}_C+1$  )
  - Server allocates space for socket only if valid  $\text{SN}_S$




# SYN Floods II: Massive flood

- E.g., BetCris.com, 2003
- Bot army to flood specific target: (DDoS)
  - **20,000** bots can generate **2Gb/sec** of SYNs (2003)
  - At web site:
    - Saturates network uplink or network router
    - Random source IP  $\Rightarrow$   
attack SYNs look the same as real SYNs
  - What to do ???

# Prolexic / Verisign Design

- Idea: only forward established TCP connections to site



# Estonia attack

- Attack types detected:
  - 115 ICMP floods, 4 TCP SYN floods
- Bandwidth:
  - 12 attacks, **70-95 Mbps for over 10 hours**
- All attack traffic was coming from outside Estonia
  - Estonia's solution: block all foreign traffic until attacks stopped

# Stronger attacks: TCP con flood

- Command bot army to:
  - Complete TCP connection to web site
  - Send short HTTP HEAD request
  - Repeat
- Will bypass SYN flood protection proxy
- ... but:
  - Attacker can no longer use random source IPs.
    - Reveals location of bot zombies
  - Proxy can now block or rate-limit bots.

# DNS DoS Attacks (Bluesecurity, 2006)

- DNS runs on UDP port 53
  - DNS entry for victim.com hosted at victim\_isp.com
- DDoS attack:
  - Flood victim\_isp.com with requests for victim.com
  - **Random source IP address** in UDP packets
- Takes out entire DNS server:
  - Bluesecurity DNS hosted at Tucows DNS server
  - DNS DDoS took out Tucows hosting many many sites
- What to do ???

# Root level DNS attacks

- Feb. 6, 2007:
  - Botnet attack on the 13 Internet DNS root servers
  - Lasted 2.5 hours
  - None crashed, but two performed badly:
    - g-root (US Department of Defense)
    - I-root (ICANN)
- Attack in Oct. 2002 took out 9 of the 13 TLD servers

# DNS DoS solutions

- Generic DDoS solutions:
  - Later on, require major changes to DNS.
- DoS resistant DNS design:
  - **CoDoNS**: [Sire, 2004]
    - **Cooperative** Domain Name System
  - P2P design for DNS system:
    - DNS nodes share the load
    - Simple update of DNS entries
    - Backwards compatible with existing DNS

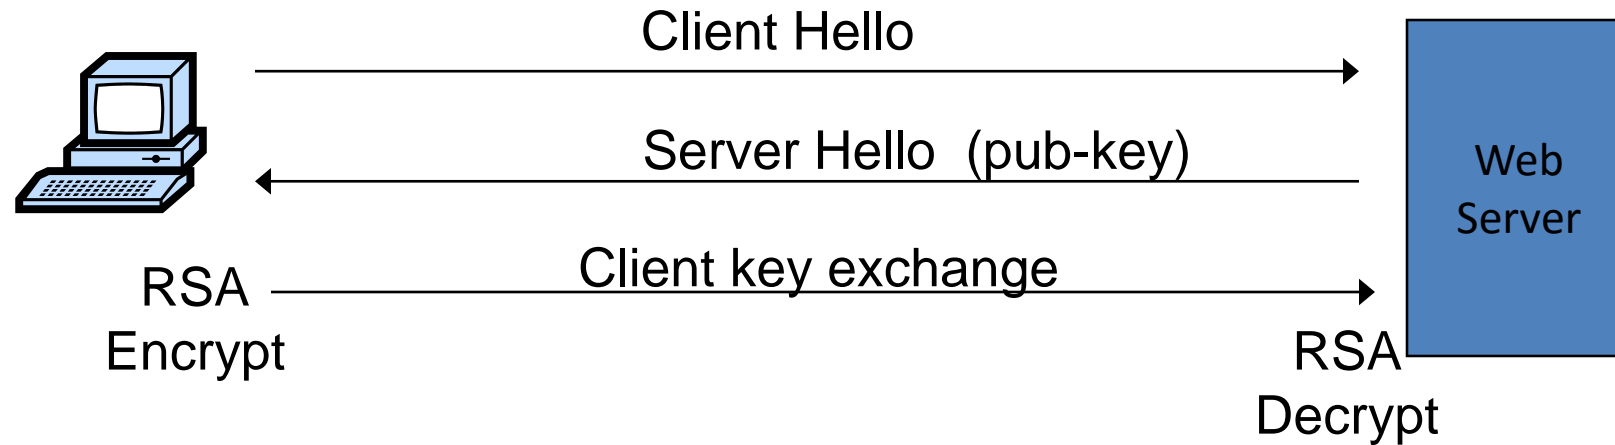
# DoS via route hijacking

- YouTube is 208.65.152.0/**22** (includes  $2^{10}$  IP addr)
  - youtube.com is 208.65.153.238, ...
- Feb. 2008:
  - Pakistan telecom advertised a BGP path for 208.65.153.0/**24** (includes  $2^8$  IP addr)
  - Routing decisions use most specific prefix
  - The entire Internet now thinks 208.65.153.238 is in Pakistan
- Outage resolved within two hours
  - ... but demonstrates huge DoS vuln. with no solution!



# DoS at higher layers

- SSL/TLS handshake [SD, 203]



RSA-encrypt speed  $\approx 10\times$  RSA-decrypt speed  
 $\Rightarrow$  Single machine can bring down ten web servers

Similar problem with application DoS:

- Send HTTP request for some large PDF file
- $\Rightarrow$  Easy work for client, hard work for server.

# **DoS Mitigation**

# 1. Client puzzles

- Idea: slow down attacker
- Moderately hard problem:
  - Given challenge  $C$  find  $X$  such that
$$\text{LSB}_n(\text{SHA-1}(C \parallel X)) = 0^n$$
  - Assumption: takes expected  $2^n$  time to solve
  - For  $n=16$  takes about 0.3sec on 1GHz machine
  - Main point: checking puzzle solution is easy
- During DoS attack:
  - Everyone must submit puzzle solution with requests
  - When no attack: do not require puzzle solution

# Examples

- TCP connection floods: (RSA, 1999)
  - Example challenge:  $C = \text{TCP server-seq-num}$
  - First data packet must contain puzzle solution
    - Otherwise TCP connection is closed
- **SSL handshake DoS:** (SD, 2003)
  - Challenge  $C$  based on TLS session ID
  - Server: check puzzle solution before RSA decrypt.
- Same for application layer DoS and payment DoS.

# Benefits and limitations

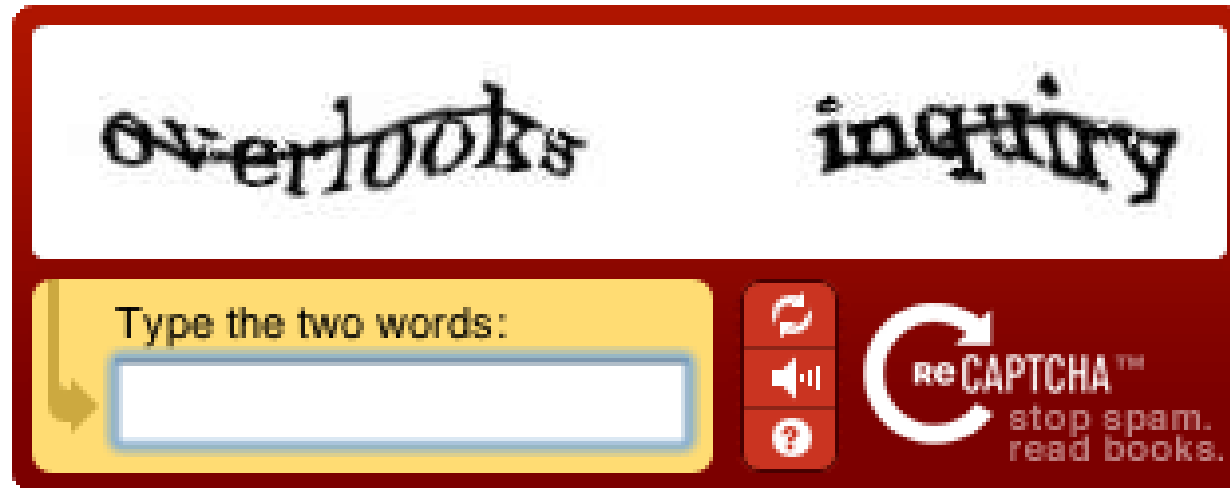
- Hardness of challenge:  $n$ 
  - Decided based on DoS attack volume.
- Limitations:
  - Requires changes to both clients and servers
  - Hurts low power legitimate clients during attack:
    - Clients on cell phones and tablets cannot connect

# Memory-bound functions

- CPU power ratio:
  - High end server / low end cell phone = 8000
  - Impossible to scale to hard puzzles
- Interesting observation:
  - Main memory access time ratio:
    - high end server / low end cell phone = 2
- Better puzzles:
  - Solution requires many main memory accesses
    - Dwork-Goldberg-Naor, Crypto 2003
    - Abadi-Burrows-Manasse-Wobber, ACM ToIT 2005

## 2. CAPTCHAs

- Idea: verify that connection is from a **human**



- Applies to application layer DDoS [Killbots 2005]
  - During attack: generate CAPTCHAs and process request only if valid solution
  - Present one CAPTCHA per source IP address.

# **3. Source identification**


Goal: identify packet source

Ultimate goal: block attack at the source



## 3a. Ingress filtering (RFC 2827, 2000)

- Big problem: DDoS with spoofed source IPs
- Question: how to find packet origin?



- Ingress filtering policy: ISP only forwards packets with legitimate source IP

# Implementation problems

- ALL ISPs must do this.      Requires global trust.
  - If 10% of ISPs do not implement  $\Rightarrow$  no defense

## 3b. Traceback [Savage et al., 2000]


- Goal:
  - Given set of attack packets, determine path to source
- How: change routers to record info in packets
- Assumptions:
  - Most routers remain uncompromised
  - Attacker sends many packets
  - Route from attacker to victim remains relatively stable

# Simple method

- Write path into network packet
  - Each router adds its own IP address to packet
  - Victim reads path from packet
- Problem:
  - Requires space in packet
    - Path can be long
    - No extra fields in current IP format
    - ... Changes to packet format too much to expect

# Better idea

- DDoS involves many packets on same path
- Store one link in each packet
  - Each router probabilistically stores own address
  - Fixed space regardless of path length




# Edge Sampling

- Data fields written to packet:
  - Edge: *start* and *end* IP addresses
  - Distance: number of hops since edge stored
- Marking procedure for router R
  - if coin turns up heads (with probability  $p$ ) then
    - write R into start address
    - write 0 into distance field
  - else
    - if distance == 0 write R into end field
    - increment distance field


# Edge Sampling: picture

- Packet received
  - $R_1$  receives packet from source or another router
  - Packet contains space for start, end, distance



# Edge Sampling: picture


- Begin writing edge
  - $R_1$  chooses to write start of edge
  - Sets distance to 0






# Edge Sampling

- Finish writing edge
  - R2 chooses not to overwrite edge
  - Distance is 0
    - Write end of edge, increment distance to 1



# Edge Sampling

- Increment distance
  - R3 chooses not to overwrite edge
  - Distance > 0
    - Increment distance to 2



# Path reconstruction

- Extract information from attack packets
- Build graph rooted at victim
  - Each (start,end,distance) tuple provides an edge
- # packets needed to reconstruct path

$$E(X) < \frac{\ln(d)}{p(1-p)^{d-1}}$$

where  $p$  is marking probability,  $d$  is length of path

# Details: where to store edge

- Identification field
  - Used for fragmentation
  - Fragmentation is rare
  - 16 bits
- Store edge in 16 bits?

offset	distance	edge chunk
0	2 3	7 8 15

- Break into chunks
- Store start + end

Version	Header Length
Type of Service	
Total Length	
Identification	
Flags	Fragment Offset
Time to Live	
Protocol	
Header Checksum	
Source Address of Originating Host	
Destination Address of Target Host	
Options	
Padding	
IP Data	

# **Capability based defense**

# Capability based defense

- Anderson, Roscoe, Wetherall.
  - Preventing internet denial-of-service with capabilities. SIGCOMM, 2004.
- Yaar, Perrig, and Song.
  - Siff: A stateless internet flow filter to mitigate DDoS flooding attacks. IEEE S&P, 2004.
- Yang, Wetherall, Anderson.
  - A DoS-limiting network architecture. SIGCOMM, 2005

# Pushback filtering

- Mahajan, Bellovin, Floyd, Ioannidis, Paxson, Shenker. Controlling High Bandwidth Aggregates in the Network. *Computer Communications Review* 2002.
- Ioannidis, Bellovin.  
Implementing Pushback: Router-Based Defense Against DoS Attacks. *NDSS* 2002
- Argyraki, Cheriton.  
Active Internet Traffic Filtering: Real-Time Response to Denial-of-Service Attacks. *USENIX* 2005.

# Overlay filtering

- Keromytis, Misra, Rubenstein.  
SOS: Secure Overlay Services. SIGCOMM 2002.
- D. Andersen. Mayday.  
Distributed Filtering for Internet Services.  
Usenix USITS 2003.
- Lakshminarayanan, Adkins, Perrig, Stoica.  
Taming IP Packet Flooding Attacks. HotNets 2003.



# Take home message:

- Denial of Service attacks are real and must be considered at design time
- Sad truth:
  - Current Internet is ill-equipped to handle DDoS attacks
- Many good proposals for core redesign...

# Conclusion

- Denial of Service
  - Attempt to defeat **availability**
    - E.g., user is denied access to service or data
- Flooding or Overload
  - Presenting commands **more quickly** than a server can handle them
    - Targeting applications or resources
- Blocked access / access failure
  - Preventing a service from functioning

# Vulnerabilities

- Internet not designed to handle DoS...
- Insufficient resources (e.g., capacity)
- Traffic redirection

More... ?